

621 Assignment3

Raghu

Oct 20, 2018

Contents

Overview:	2
1. DATA EXPLORATION	3
Basic Stats	4
Data Visualization	5
BoxPlot Distribution	5
BoxPlot of all variables	6
BoxPlot with Log Scale	7
Density Plots	8
Scatterplot	9
Correlations	10
2. DATA PREPARATION	12
ScatterPlot of log transformations	14
3. BUILD MODELS	15
Leaps Subsetting of Untransformed Data	15
Model 1: All Variables	18
Transformed Data Analysis.	19
Model 2: Transformed Variables	22
Model 3: (Logarithmic Model) with stepwise.	23
Model 4: (Logarithmic Model) with Principal Components.	26
4. SELECT MODELS	28
Test Model1	30
Performance	30
Prediction for Test Data	32
Appendix	33

Overview:

In this homework assignment, you will explore, analyze and model a data set containing information on crime for various neighborhoods of a major city. Each record has a response variable indicating whether or not the crime rate is above the median crime rate (1) or not (0).

Your objective is to build a binary logistic regression model on the training data set to predict whether the neighborhood will be at risk for high crime levels. You will provide classifications and probabilities for the evaluation data set using your binary logistic regression model. You can only use the variables given to you (or variables that you derive from the variables provided). Below is a short description of the variables of interest in the data set:

Explanatory Variables:

zn: proportion of residential land zoned for large lots (over 25000 square feet) (predictor variable)

indus: proportion of non-retail business acres per suburb (predictor variable)

chas: a dummy var. for whether the suburb borders the Charles River (1) or not (0) (predictor variable)

nox: nitrogen oxides concentration (parts per 10 million) (predictor variable)

rm: average number of rooms per dwelling (predictor variable)

age: proportion of owner-occupied units built prior to 1940 (predictor variable)

dis: weighted mean of distances to five Boston employment centers (predictor variable)

rad: index of accessibility to radial highways (predictor variable)

tax: full-value property-tax rate per \$10,000 (predictor variable)

ptratio: pupil-teacher ratio by town (predictor variable)

black: $1000(Bk - 0.63)^2$ where Bk is the proportion of blacks by town (predictor variable)

lstat: lower status of the population (percent) (predictor variable)

medv: median value of owner-occupied homes in \$1000s (predictor variable)

target: whether the crime rate is above the median crime rate (1) or not (0) (response variable)

1. DATA EXPLORATION

Describe the size and the variables in the crime training data set. Consider that too much detail will cause a manager to lose interest while too little detail will make the manager consider that you aren't doing your job.

- Mean / Standard Deviation / Median
- Bar Chart or Box Plot of the data
- Is the data correlated to the target variable (or to other variables?)
- Are any of the variables missing and need to be imputed "fixed"?

Data View:

Lets have a quick view of crime data.

zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	black	lstat	medv	target
0	19.58	0	0.605	7.929	96.2	2.0459	5	403	14.7	369.30	3.70	50.0	1
0	19.58	1	0.871	5.403	100.0	1.3216	5	403	14.7	396.90	26.82	13.4	1
0	18.10	0	0.740	6.485	100.0	1.9784	24	666	20.2	386.73	18.85	15.4	1
30	4.93	0	0.428	6.393	7.8	7.0355	6	300	16.6	374.71	5.19	23.7	0
0	2.46	0	0.488	7.155	92.2	2.7006	3	193	17.8	394.12	4.82	37.9	0
0	8.56	0	0.520	6.781	71.3	2.8561	5	384	20.9	395.58	7.67	26.5	0

Basic Stats

There are 466 observations and 14 variables. * 10 variables of type dbl. * 4 variables of type int.

```
## Observations: 466
## Variables: 14
## $ zn      <dbl> 0, 0, 0, 30, 0, 0, 0, 0, 0, 80, 22, 0, 0, 22, 0, 0, 10...
## $ indus   <dbl> 19.58, 19.58, 18.10, 4.93, 2.46, 8.56, 18.10, 18.10, 5...
## $ chas    <int> 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...
## $ nox     <dbl> 0.605, 0.871, 0.740, 0.428, 0.488, 0.520, 0.693, 0.693...
## $ rm      <dbl> 7.929, 5.403, 6.485, 6.393, 7.155, 6.781, 5.453, 4.519...
## $ age     <dbl> 96.2, 100.0, 100.0, 7.8, 92.2, 71.3, 100.0, 100.0, 38...
## $ dis     <dbl> 2.0459, 1.3216, 1.9784, 7.0355, 2.7006, 2.8561, 1.4896...
## $ rad     <int> 5, 5, 24, 6, 3, 5, 24, 24, 5, 1, 7, 5, 24, 7, 3, 3, 5,...
## $ tax     <int> 403, 403, 666, 300, 193, 384, 666, 666, 224, 315, 330,...
## $ ptratio <dbl> 14.7, 14.7, 20.2, 16.6, 17.8, 20.9, 20.2, 20.2, 20.2, ...
## $ black   <dbl> 369.30, 396.90, 386.73, 374.71, 394.12, 395.58, 396.90...
## $ lstat   <dbl> 3.70, 26.82, 18.85, 5.19, 4.82, 7.67, 30.59, 36.98, 5...
## $ medv    <dbl> 50.0, 13.4, 15.4, 23.7, 37.9, 26.5, 5.0, 7.0, 22.2, 20...
## $ target  <int> 1, 1, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, ...
```

- Summary Statistics shows none of the variables have missing values
- The mean of target is below 0.5 which means there are more observations where the crime rate is below the median.

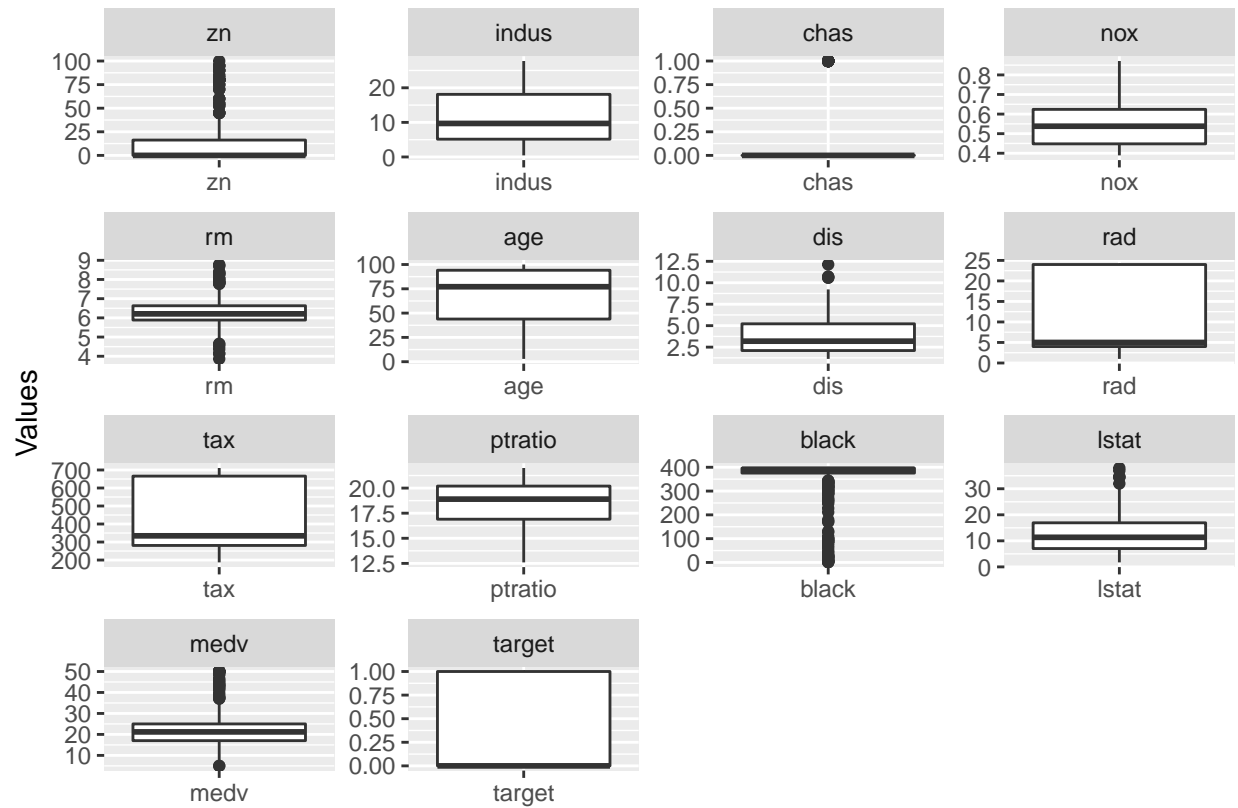
	vars	n	mean	sd	median	trimmed	mad	min	max	ra
zn	1	466	11.5772532	23.3646511	0.00000	5.3542781	0.0000000	0.0000	100.0000	100.0
indus	2	466	11.1050215	6.8458549	9.69000	10.9082353	9.3403800	0.4600	27.7400	27.2
chas	3	466	0.0708155	0.2567920	0.00000	0.0000000	0.0000000	0.0000	1.0000	1.0
nox	4	466	0.5543105	0.1166667	0.53800	0.5442684	0.1334340	0.3890	0.8710	0.4
rm	5	466	6.2906738	0.7048513	6.21000	6.2570615	0.5166861	3.8630	8.7800	4.9
age	6	466	68.3675966	28.3213784	77.15000	70.9553476	30.0226500	2.9000	100.0000	97.1
dis	7	466	3.7956929	2.1069496	3.19095	3.5443647	1.9144814	1.1296	12.1265	10.9
rad	8	466	9.5300429	8.6859272	5.00000	8.6978610	1.4826000	1.0000	24.0000	23.0
tax	9	466	409.5021459	167.9000887	334.50000	401.5080214	104.5233000	187.0000	711.0000	524.0
ptratio	10	466	18.3984979	2.1968447	18.90000	18.5970588	1.9273800	12.6000	22.0000	9.4
black	11	466	357.1201502	91.3211298	391.34000	383.5064439	8.2432560	0.3200	396.9000	396.5
lstat	12	466	12.6314592	7.1018907	11.35000	11.8809626	7.0720020	1.7300	37.9700	36.2
medv	13	466	22.5892704	9.2396814	21.20000	21.6304813	6.0045300	5.0000	50.0000	45.0
target	14	466	0.4914163	0.5004636	0.00000	0.4893048	0.0000000	0.0000	1.0000	1.0

Data Visualization


BoxPlot Distribution

Boxplot demonstrating the mean, median and quartiles of the independent variables. rad, tax and black has high variances.

No id variables; using all as measure variables



BoxPlot of all variables




The box plots displays that many of the variables have low variances.

Lets look at the log scale of the independent variables.

BoxPlot with Log Scale

```
## Warning: Transformation introduced infinite values in continuous y-axis
```

```
## Warning: Removed 1009 rows containing non-finite values (stat_boxplot).
```



Density Plots

Lets look at the Density Plots for skewness.

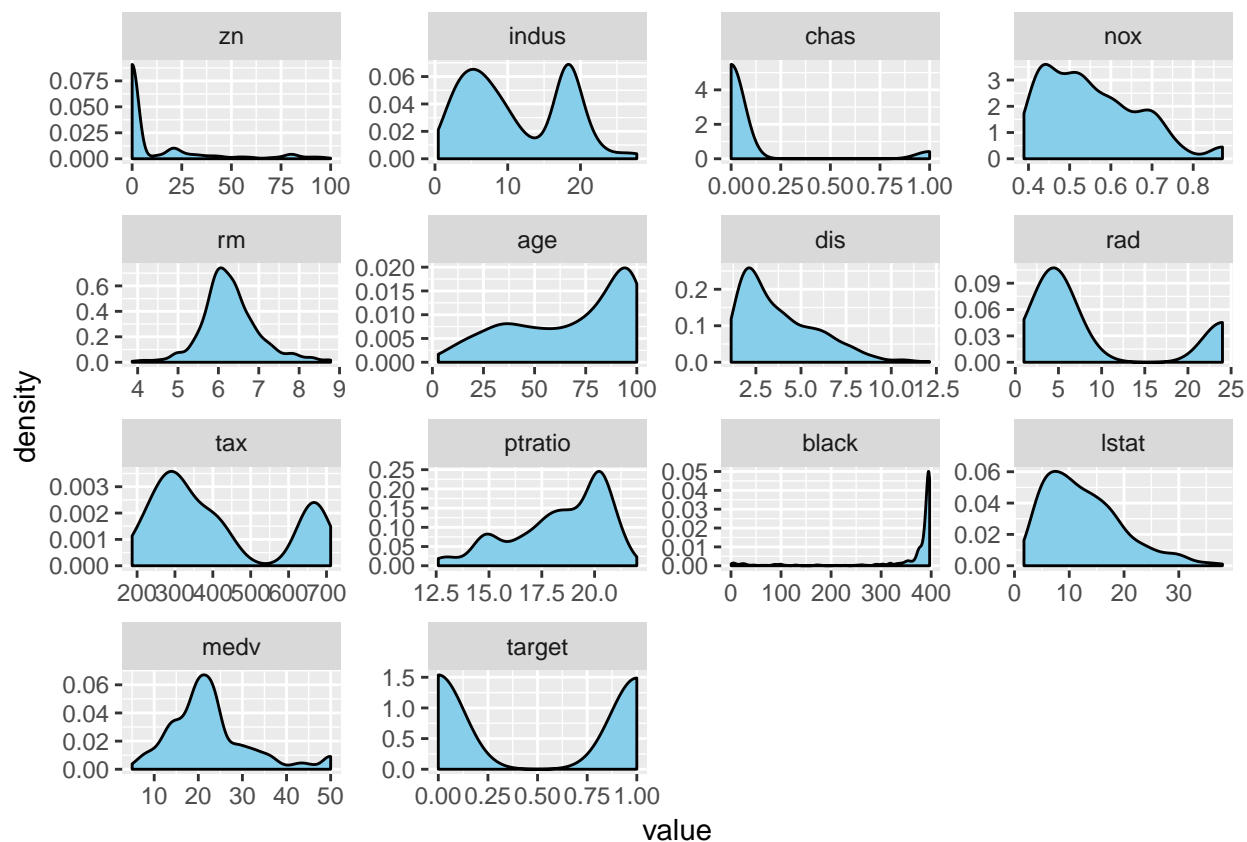
-rm is the only variable that closely mirrors a normal distribution.

-zn, chas, and dis are heavily skewed right.

-nox, lstat, and medv are are also skewed right.


-indus, rad, tax, and target are multi-modal.

The density plots reveal that most is the data is not normal.




Scatterplot

Interpreting a binomial response variable may not be the most best way to visualize the data using scatterplot.



Correlations

“nox” has the highest positive correlation and “dis” has the highest negative correlation. “tax” and “rad” are correlated to each other.



```
## [1] "Correlation to Target"
```

	target
zn	-0.4316818
indus	0.6048507
chas	0.0800419
nox	0.7261062
rm	-0.1525533
age	0.6301062
dis	-0.6186731
rad	0.6281049
tax	0.6111133
ptratio	0.2508489
black	-0.3529568
lstat	0.4691270
medv	-0.2705507
target	1.0000000

```
## [1] "Positive Correlative Factors:"
```

```
## [1] "indus" "nox" "age" "rad" "tax" "ptratio" "lstat"
## [8] "target"
```

```
## [1] "Negative Correlative Factors:"
## [1] "zn"      "dis"      "black" "medv"
## [1] "Neutral Correlative Factors"
## [1] "chas" "rm"
## [1] "Highly Positive Correlated Variables"
```

	target
nox	0.7261062

```
## [1] "Highly Negative Correlated Variables"
```

	target
dis	-0.6186731

2. DATA PREPARATION

Describe how you have transformed the data by changing the original variables or creating new variables. If you did transform the data or create new variables, discuss why you did this. Here are some possible transformations.

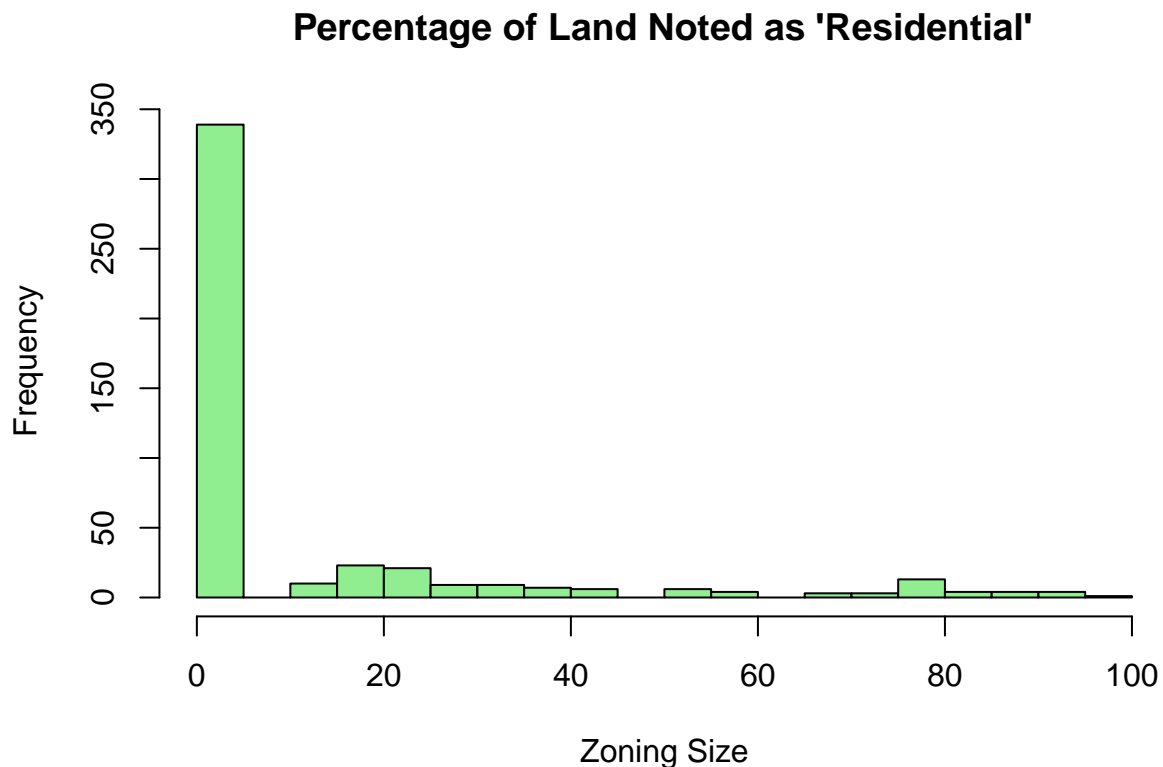
- Fix missing values (maybe with a Mean or Median value)
- Create flags to suggest if a variable was missing
- Transform data by putting it into buckets
- Mathematical transforms such as log or square root (or use Box-Cox)
- Combine variables (such as ratios or adding or multiplying) to create new variables

Transformation:

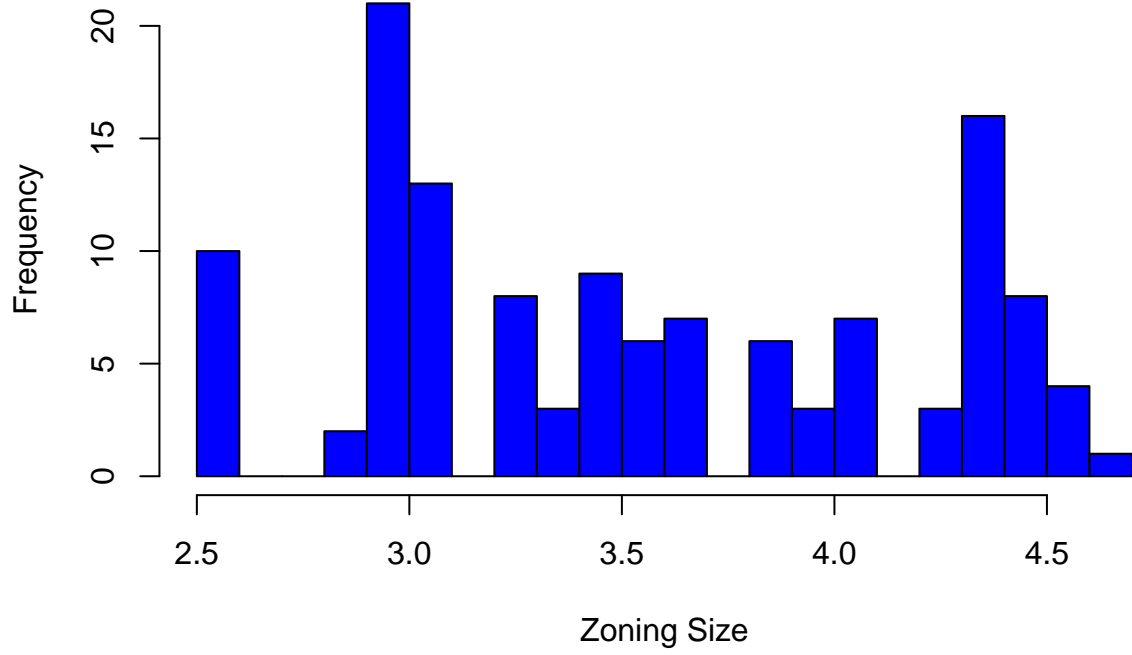
To reduce the effect of skewness on the model, lets do log transformations on all the variables except the variables that are binary(Zn,chas).

Due to high correlation between two indepdent variables that is between tax and rad, We can build an interactive term for this when we build our models as they are possibly likely very dependent on each other with one term affecting the other.

Lets look at Zn:proportion of residential land zoned for large lots



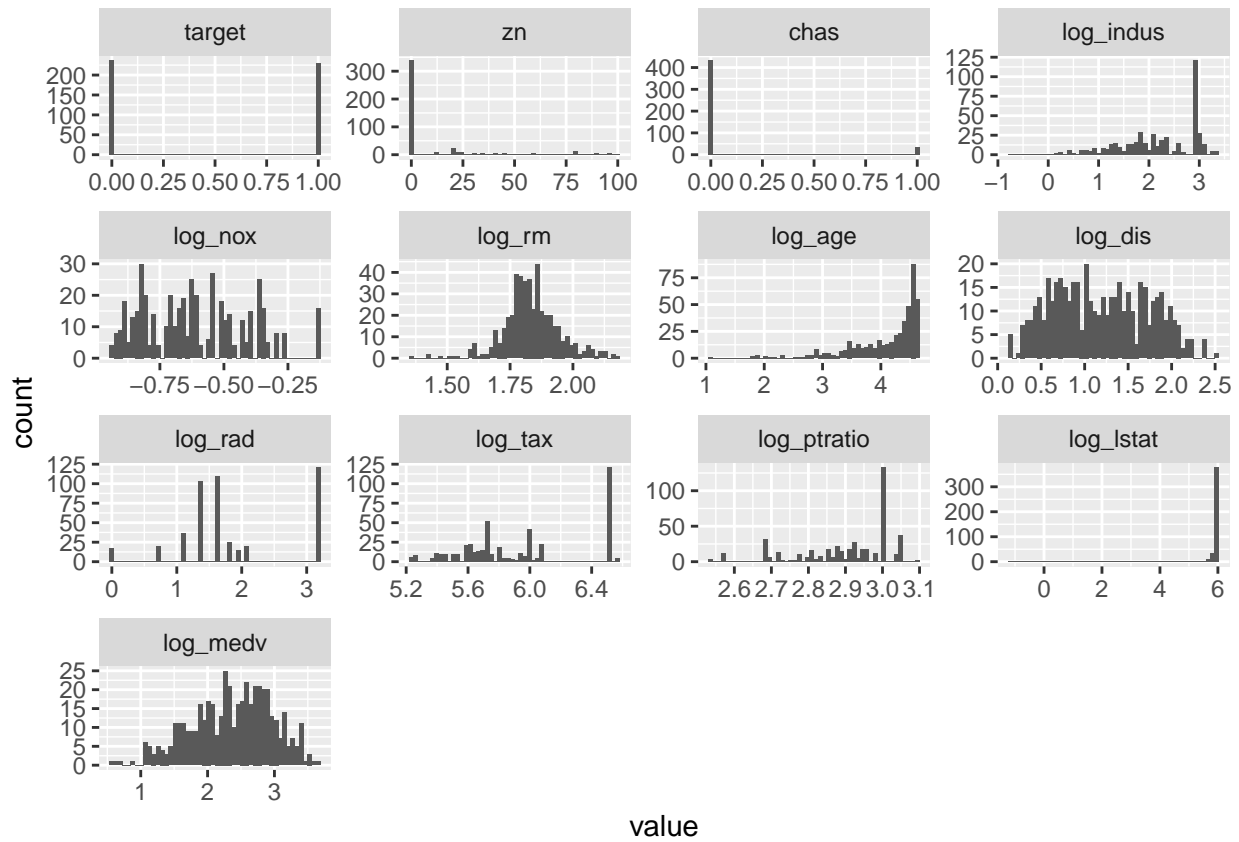
Percentage of Logarithmic Land Noted as 'Residential'



Let's create a scatterplot with these new variables for the logarithmic transformations.

ScatterPlot of log transformations

No id variables; using all as measure variables



3. BUILD MODELS

Using the training data, build at least three different binary logistic regression models, using different variables (or the same variables with different transformations). You may select the variables manually, use an approach such as Forward or Stepwise, use a different approach, or use a combination of techniques. Describe the techniques you used. If you manually selected a variable for inclusion into the model or exclusion into the model, indicate why this was done.


Be sure to explain how you can make inferences from the model, as well as discuss other relevant model output. Discuss the coefficients in the models, do they make sense? Are you keeping the model even though it is counter intuitive? Why? The boss needs to know.

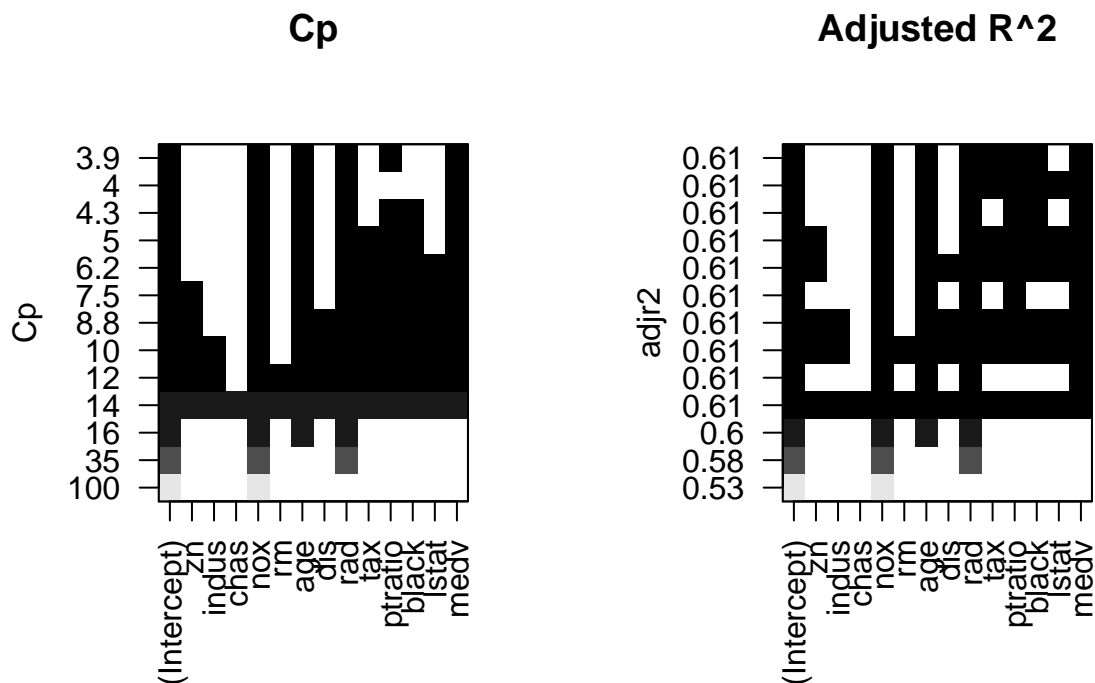
Leaps Subsetting of Untransformed Data

The Leaps package is an “regression subset selection” tool. The package automatically generates all possible models. The tool is basically used to find the “best” model.

```
## Subset selection object
## Call: regsubsets.formula(target ~ ., data = crime_train, method = "exhaustive",
##      nvmax = NULL, nbest = 1)
## 13 Variables (and intercept)
##      Forced in Forced out
## zn          FALSE      FALSE
## indus        FALSE      FALSE
## chas         FALSE      FALSE
## nox          FALSE      FALSE
## rm           FALSE      FALSE
## age          FALSE      FALSE
## dis          FALSE      FALSE
## rad          FALSE      FALSE
## tax          FALSE      FALSE
## ptratio      FALSE      FALSE
## black        FALSE      FALSE
## lstat        FALSE      FALSE
## medv         FALSE      FALSE
## 1 subsets of each size up to 13
## Selection Algorithm: exhaustive
##      zn  indus chas nox rm  age dis rad tax ptratio black lstat medv
## 1 ( 1 ) " " " " " " "*" " " " " " " " " " " " " " " " "
## 2 ( 1 ) " " " " " " "*" " " " " " " "*" " " " " " " " "
## 3 ( 1 ) " " " " " " "*" " " " "*" " " "*" " " " " " " " "
## 4 ( 1 ) " " " " " " "*" " " " "*" " " "*" " " " " " " "*"
## 5 ( 1 ) " " " " " " "*" " " " "*" " " "*" " " " "*" " " "*"
## 6 ( 1 ) " " " " " " "*" " " " "*" " " "*" " " " "*" " " "*"
## 7 ( 1 ) " " " " " " "*" " " " "*" " " "*" "*" "*" " "*" " "*"
## 8 ( 1 ) " " " " " " "*" " " " "*" " " "*" "*" "*" " "*" "*"
## 9 ( 1 ) "*" " " " " " "*" " " " "*" " " "*" "*" "*" " "*" "*"
## 10 ( 1 ) "*" " " " " " "*" " " " "*" "*" "*" "*" " "*" "*"
## 11 ( 1 ) "*" "*" " " " " "*" " " " "*" "*" "*" "*" " "*" "*"
## 12 ( 1 ) "*" "*" " " " " "*" "*" "*" "*" "*" "*" "*" " "*" "*"
## 13 ( 1 ) "*" "*" "*" "*" "*" "*" "*" "*" "*" "*" "*" " "*" "*"

```





Based on Cp, a model that includes nox, age, rad, ptratio, and medv would be the best predictor.

Based on Adjusted R², a model that includes nox, age, rad, tax, ptratio, black, and medv would be the best predictor.

Both metrics share the nox, age, rad, ptratio, and medv variables.

Model 1: All Variables

The glmulti package is an “automated model selection and model averaging” tool. The package automatically generates all possible models “with the specified response and explanatory variables”.

All of the variables will be tested to determine the base model they provided. This will allow us to see which variables are significant in our dataset, and allow us to make other models based on that. This model will be based off of the original data - before transformed (log) variables have been added to account for potential issues in the data.

“nox”, “rad”, “ptratio” are highly statistically significant and “dis” and “medv” are somewhat significant. “nox” has high impact on target. tax has minimum impact and also negative.

Positive coefficients: chas, nox, age, dis, rad, ptratio, lstat, medv

Negative coefficients: zn, indus, rm, tax, black


lets see how the other models reports on the deviance and AIC for comparison.


```
##
## Call:
## glm(formula = target ~ ., family = binomial(link = "logit"),
##      data = crime_train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.2854  -0.1372  -0.0017   0.0020   3.4721
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -36.839521   7.028726  -5.241 1.59e-07 ***
## zn          -0.061720   0.034410  -1.794 0.072868 .
## indus       -0.072580   0.048546  -1.495 0.134894
## chas         1.032352   0.759627   1.359 0.174139
## nox         50.159513   8.049503   6.231 4.62e-10 ***
## rm          -0.692145   0.741431  -0.934 0.350548
## age          0.034522   0.013883   2.487 0.012895 *
## dis          0.765795   0.234407   3.267 0.001087 **
## rad          0.663015   0.165135   4.015 5.94e-05 ***
## tax         -0.006593   0.003064  -2.152 0.031422 *
## ptratio      0.442217   0.132234   3.344 0.000825 ***
## black       -0.013094   0.006680  -1.960 0.049974 *
## lstat        0.047571   0.054508   0.873 0.382802
## medv         0.199734   0.071022   2.812 0.004919 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 645.88  on 465  degrees of freedom
## Residual deviance: 186.15  on 452  degrees of freedom
## AIC: 214.15
##
## Number of Fisher Scoring iterations: 9
```

Transformed Data Analysis.

Lets look at the transformed Data.

```
## Subset selection object
## Call: regsubsets.formula(target ~ ., data = crime_train_log, method = "exhaustive",
##      nvmax = NULL, nbest = 1)
## 12 Variables (and intercept)
##      Forced in Forced out
## zn            FALSE      FALSE
## chas          FALSE      FALSE
## log_indus     FALSE      FALSE
## log_nox       FALSE      FALSE
## log_rm        FALSE      FALSE
## log_age       FALSE      FALSE
## log_dis       FALSE      FALSE
## log_rad       FALSE      FALSE
## log_tax       FALSE      FALSE
## log_ptratio   FALSE      FALSE
## log_lstat     FALSE      FALSE
## log_medv      FALSE      FALSE
## 1 subsets of each size up to 12
## Selection Algorithm: exhaustive
##      zn chas log_indus log_nox log_rm log_age log_dis log_rad
## 1 ( 1 ) " " " " " " " " " " " "
## 2 ( 1 ) " " " " " " " " " " " "
## 3 ( 1 ) " " " " " " " " " " " "
## 4 ( 1 ) " " " " " " " " " " " "
## 5 ( 1 ) " " " " " " " " " " " "
## 6 ( 1 ) " " " " " " " " " " " "
## 7 ( 1 ) " " " " " " " " " " " "
## 8 ( 1 ) " " " " " " " " " " " "
## 9 ( 1 ) " " " " " " " " " " " "
## 10 ( 1 ) " " " " " " " " " " " "
## 11 ( 1 ) " " " " " " " " " " " "
## 12 ( 1 ) " " " " " " " " " " " "
##      log_tax log_ptratio log_lstat log_medv
## 1 ( 1 ) " " " " " "
## 2 ( 1 ) " " " " " "
## 3 ( 1 ) " " " " " "
## 4 ( 1 ) " " " " " "
## 5 ( 1 ) " " " " " "
## 6 ( 1 ) " " " " " "
## 7 ( 1 ) " " " " " "
## 8 ( 1 ) " " " " " "
## 9 ( 1 ) " " " " " "
## 10 ( 1 ) " " " " " "
## 11 ( 1 ) " " " " " "
## 12 ( 1 ) " " " " " "
```





CP has reduced to 4 from 5 after transformation.

Both CP and Rsquare indicates “nox”, “rm” , age“,”rad" are best predictors.

Model 2: Transformed Variables

Model2 is the log transformation of all the variables and the interactive term is included.

The log variables should help negate the large amount of skew in the data - or help them to become more normalized.

```
##
## Call:
## glm(formula = target ~ . + log_rad:log_tax, family = binomial(link = "logit"),
##      data = crime_train_log)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.00318  -0.17093  -0.00164   0.10619   3.13261
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -2.46216    14.45018  -0.170  0.86470
## zn            -0.03409     0.02674  -1.275  0.20227
## chas           0.94719     0.77212   1.227  0.21992
## log_indus      0.34000     0.56795   0.599  0.54941
## log_nox       22.81455     3.76435   6.061 1.36e-09 ***
## log_rm         5.30404     2.99089   1.773  0.07616 .
## log_age        0.48567     0.56896   0.854  0.39332
## log_dis        2.13517     0.78734   2.712  0.00669 **
## log_rad        6.89690     7.26768   0.949  0.34263
## log_tax       -1.33460     1.69875  -0.786  0.43208
## log_ptratio    3.97363     1.82156   2.181  0.02915 *
## log_lstat     -1.17666     1.13585  -1.036  0.30023
## log_medv      -0.16766     0.61577  -0.272  0.78541
## log_rad:log_tax -0.63883     1.18935  -0.537  0.59118
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 645.88  on 465  degrees of freedom
## Residual deviance: 207.08  on 452  degrees of freedom
## AIC: 235.08
##
## Number of Fisher Scoring iterations: 8
```

Analysis:

nox has the greatest impact on target.

nox, rad are highly statistically significant.

AIC has increased compared to model1.

Null deviance is the same. Residual deviance has increased.

interactive term did not add much value.

So, this model may not be the best choice. let me try stepwise for Model1 and Model2.

Model 3: (Logarithmic Model) with stepwise.

Let me try both forward and backward elimination stepwise algorithm here.

```
## Start:  AIC=235.08
## target ~ zn + chas + log_indus + log_nox + log_rm + log_age +
##      log_dis + log_rad + log_tax + log_ptratio + log_lstat + log_medv +
##      log_rad:log_tax
##
##           Df Deviance    AIC
## - log_medv      1   207.15 233.15
## - log_rad:log_tax 1   207.35 233.35
## - log_indus      1   207.44 233.44
## - log_age        1   207.84 233.84
## - log_lstat      1   208.53 234.53
## - chas           1   208.61 234.61
## - zn             1   208.99 234.99
## <none>           207.08 235.08
## - log_rm         1   210.17 236.17
## - log_ptratio    1   211.92 237.92
## - log_dis        1   214.86 240.86
## - log_nox        1   267.00 293.00
##
## Step:  AIC=233.15
## target ~ zn + chas + log_indus + log_nox + log_rm + log_age +
##      log_dis + log_rad + log_tax + log_ptratio + log_lstat + log_rad:log_tax
##
##           Df Deviance    AIC
## - log_rad:log_tax 1   207.42 231.42
## - log_indus      1   207.50 231.50
## - log_age        1   207.85 231.85
## - log_lstat      1   208.61 232.61
## - chas           1   208.63 232.63
## <none>           207.15 233.15
## - zn             1   209.24 233.24
## + log_medv       1   207.08 235.08
## - log_ptratio    1   211.93 235.93
## - log_rm         1   214.62 238.62
## - log_dis        1   214.88 238.88
## - log_nox        1   267.25 291.25
##
## Step:  AIC=231.42
## target ~ zn + chas + log_indus + log_nox + log_rm + log_age +
##      log_dis + log_rad + log_tax + log_ptratio + log_lstat
##
##           Df Deviance    AIC
## - log_indus      1   207.62 229.62
## - log_age        1   208.11 230.11
## - log_lstat      1   208.75 230.75
## - chas           1   209.24 231.24
## <none>           207.42 231.42
## - zn             1   209.57 231.57
## + log_rad:log_tax 1   207.15 233.15
## - log_tax        1   211.31 233.31
```

```

## + log_medv      1   207.35 233.35
## - log_ptratio   1   212.07 234.07
## - log_rm        1   215.38 237.38
## - log_dis       1   216.15 238.15
## - log_rad       1   248.83 270.83
## - log_nox       1   268.82 290.82
##
## Step:  AIC=229.62
## target ~ zn + chas + log_nox + log_rm + log_age + log_dis + log_rad +
##         log_tax + log_ptratio + log_lstat
##
##           Df Deviance   AIC
## - log_age      1   208.28 228.28
## - log_lstat     1   208.99 228.99
## <none>          207.62 229.62
## - chas         1   209.88 229.88
## - zn           1   210.12 230.12
## + log_indus     1   207.42 231.42
## + log_rad:log_tax 1   207.50 231.50
## + log_medv      1   207.56 231.56
## - log_tax       1   211.64 231.64
## - log_ptratio   1   212.53 232.53
## - log_rm        1   215.41 235.41
## - log_dis       1   216.40 236.40
## - log_rad       1   249.58 269.58
## - log_nox       1   278.93 298.93
##
## Step:  AIC=228.28
## target ~ zn + chas + log_nox + log_rm + log_dis + log_rad + log_tax +
##         log_ptratio + log_lstat
##
##           Df Deviance   AIC
## - log_lstat     1   209.76 227.76
## <none>          208.28 228.28
## - chas         1   210.87 228.87
## - zn           1   211.06 229.06
## + log_age       1   207.62 229.62
## + log_indus     1   208.11 230.11
## - log_tax       1   212.12 230.12
## + log_rad:log_tax 1   208.17 230.17
## + log_medv      1   208.27 230.27
## - log_ptratio   1   213.08 231.08
## - log_rm        1   216.16 234.16
## - log_dis       1   216.50 234.50
## - log_rad       1   249.61 267.61
## - log_nox       1   287.61 305.61
##
## Step:  AIC=227.76
## target ~ zn + chas + log_nox + log_rm + log_dis + log_rad + log_tax +
##         log_ptratio
##
##           Df Deviance   AIC
## <none>          209.76 227.76
## - chas         1   212.20 228.20

```



```

## + log_lstat      1   208.28 228.28
## - zn             1   212.64 228.64
## + log_age        1   208.99 228.99
## - log_tax        1   213.37 229.37
## + log_indus      1   209.55 229.55
## + log_rad:log_tax 1   209.73 229.73
## + log_medv       1   209.74 229.74
## - log_ptratio    1   214.75 230.75
## - log_rm         1   217.34 233.34
## - log_dis        1   217.64 233.64
## - log_rad        1   255.17 271.17
## - log_nox        1   290.84 306.84

##
## Call:
## glm(formula = target ~ zn + chas + log_nox + log_rm + log_dis +
##      log_rad + log_tax + log_ptratio, family = binomial(link = "logit"),
##      data = crime_train_log)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.93386  -0.18898  -0.00247   0.09744   3.11267
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -3.89208     8.50307  -0.458  0.64715
## zn           -0.04000     0.02603  -1.537  0.12439
## chas          1.10355     0.72056   1.532  0.12565
## log_nox       23.90164     3.52542   6.780 1.20e-11 ***
## log_rm        5.68049     2.10479   2.699  0.00696 **
## log_dis       2.05974     0.75548   2.726  0.00640 **
## log_rad       2.97131     0.60525   4.909 9.14e-07 ***
## log_tax      -1.75087     0.93090  -1.881  0.05999 .
## log_ptratio   3.86798     1.74937   2.211  0.02703 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 645.88  on 465  degrees of freedom
## Residual deviance: 209.76  on 457  degrees of freedom
## AIC: 227.76
##
## Number of Fisher Scoring iterations: 8

```

Analysis:

AIC has decreased compared to model-2 but still above model-1.

No difference on the Null deviance and Residual deviance.

“nox”, “rad” are highly significant and “rm”, “dis” are less significant.

Model 4: (Logarithmic Model) with Principal Components.

want to check how many variables are selected in this model.

```
## Warning in train.default(x, y, weights = w, ...): You are trying to do
## regression and your outcome only has two possible values Are you trying to
## do classification? If so, use a 2 level factor as your outcome column.
```


Model summary

```
## Data:      X dimension: 466 12
## Y dimension: 466 1
## Fit method: svdpc
## Number of components considered: 6
## TRAINING: % variance explained
##           1 comps 2 comps 3 comps 4 comps 5 comps 6 comps
## X           47.34   58.58   68.41   75.58   82.02   87.88
## .outcome    50.27   54.85   57.81   57.99   58.42   59.37
```

Model Results

ncomp	RMSE	Rsquared	MAE	RMSESD	RsquaredSD	MAESD
1	0.3524385	0.5178057	0.2812633	0.0417900	0.1121874	0.0387359
2	0.3366375	0.5565167	0.2663780	0.0415044	0.1106530	0.0362781
3	0.3259484	0.5781348	0.2554683	0.0345522	0.0888443	0.0321500
4	0.3261228	0.5776264	0.2545356	0.0325305	0.0852492	0.0312235
5	0.3249427	0.5793505	0.2542360	0.0330624	0.0855554	0.0316626
6	0.3207876	0.5900258	0.2403649	0.0333483	0.0860768	0.0331914
7	0.3209674	0.5893919	0.2400574	0.0335698	0.0867198	0.0330836
8	0.3214561	0.5886781	0.2418056	0.0329510	0.0855423	0.0320599
9	0.3213672	0.5895202	0.2414771	0.0325073	0.0829208	0.0312326
10	0.3209735	0.5898968	0.2425102	0.0318672	0.0812026	0.0299609

Model Plot



	ncomp
6	6

Analysis:

This model has selected upto 6 components.

4. SELECT MODELS

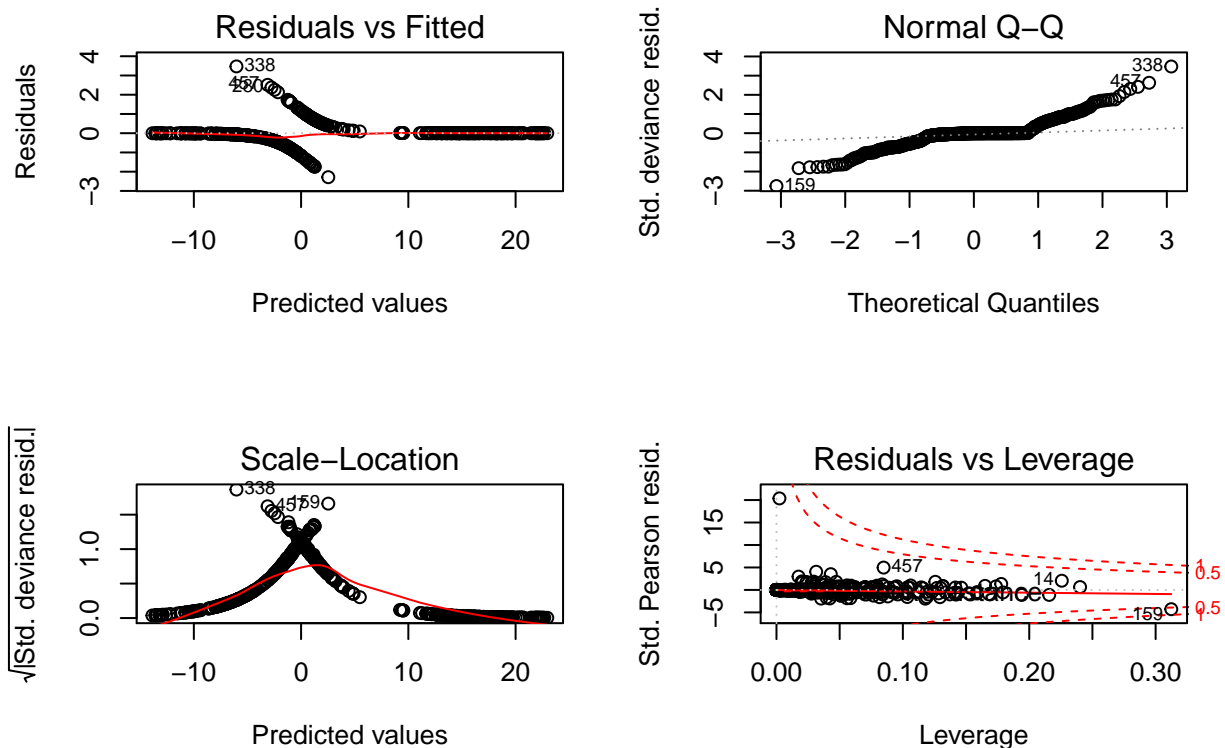
Decide on the criteria for selecting the best binary logistic regression model. Will you select models with slightly worse performance if it makes more sense or is more parsimonious? Discuss why you selected your models.

For the binary logistic regression model, will you use a metric such as log likelihood, AIC, ROC curve, etc.? Using the training data set, evaluate the binary logistic regression model based on (a) accuracy, (b) classification error rate, (c) precision, (d) sensitivity, (e) specificity, (f) F1 score, (g) AUC, and (h) confusion matrix. Make predictions using the evaluation data set.


Analysis:


After looking at the residual deviance scores and AIC scores in the previous section, we'll evaluate the model here.

Let us evaluate the Model Number 1 (baseline model). Next, we will develop a confusion matrix and create our evaluations there.



Histogram of m1\$residuals





The histogram of the residuals do not show a normal distribution.

The qqplot shows a fairly linear relationship, except towards the tail end of the residuals.

The residual indicates that there is not constant variance throughout, as there is a noticable pattern around 0.

Test Model1

zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	black	lstat	medv	target	scored_target
0	19.58	0	0.605	7.929	96.2	2.0459	5	403	14.7	369.30	3.70	50.0	1	
0	19.58	1	0.871	5.403	100.0	1.3216	5	403	14.7	396.90	26.82	13.4	1	
0	18.10	0	0.740	6.485	100.0	1.9784	24	666	20.2	386.73	18.85	15.4	1	
30	4.93	0	0.428	6.393	7.8	7.0355	6	300	16.6	374.71	5.19	23.7	0	
0	2.46	0	0.488	7.155	92.2	2.7006	3	193	17.8	394.12	4.82	37.9	0	

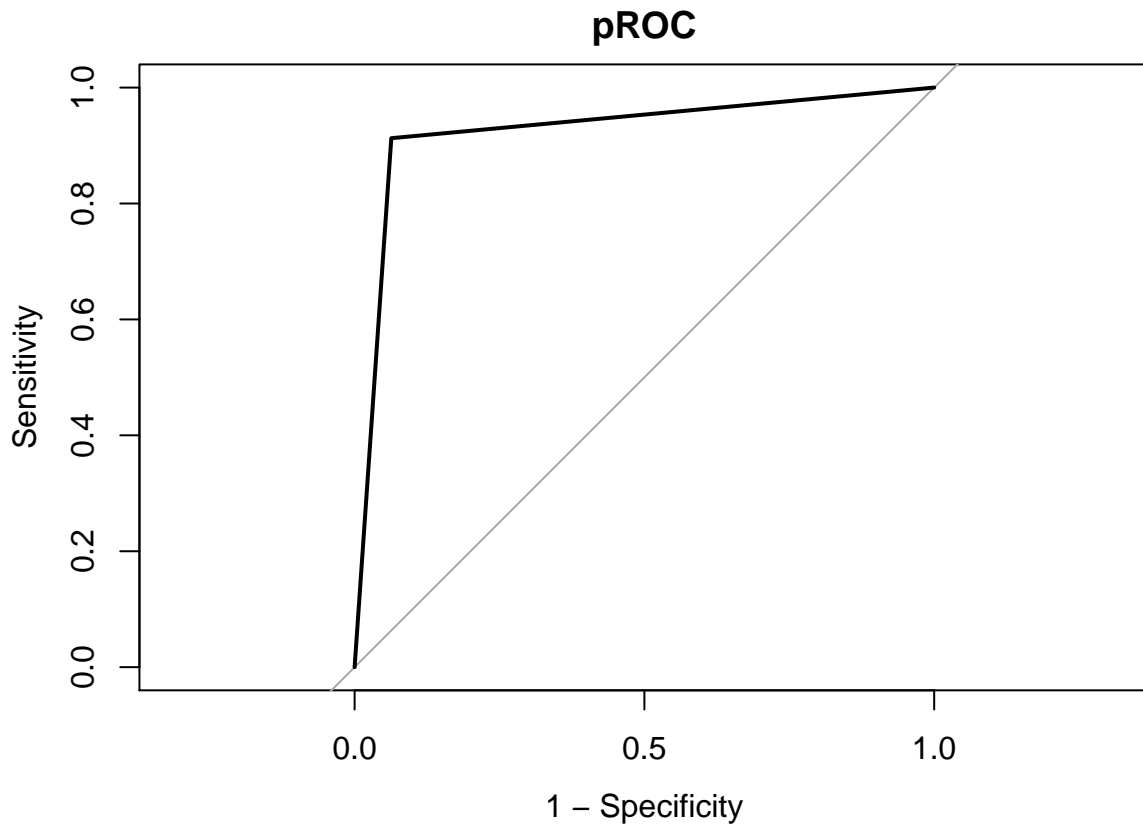
Performance

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0    1
##           0 222  20
##           1  15 209
##
##           Accuracy : 0.9249
```

```

##          95% CI : (0.8971, 0.9471)
##    No Information Rate : 0.5086
##    P-Value [Acc > NIR] : <2e-16
##
##          Kappa : 0.8497
##    McNemar's Test P-Value : 0.499
##
##          Sensitivity : 0.9127
##          Specificity : 0.9367
##          Pos Pred Value : 0.9330
##          Neg Pred Value : 0.9174
##          Precision : 0.9330
##          Recall : 0.9127
##          F1 : 0.9227
##          Prevalence : 0.4914
##          Detection Rate : 0.4485
##          Detection Prevalence : 0.4807
##          Balanced Accuracy : 0.9247
##
##          'Positive' Class : 1
##

```



Analysis:

This model has 90% accuracy. Precision is 95%. Negative prediction rate is only 91%. Positive prediction rate is 93. Sensitivity is 91% Specificity is 93% F1 is 92% AUC is 92%

Prediction for Test Data

##	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	black	lstat
## 1	0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	4.03
## 2	0	8.14	0	0.538	6.096	84.5	4.4619	4	307	21.0	380.02	10.26
## 3	0	8.14	0	0.538	6.495	94.4	4.4547	4	307	21.0	387.94	12.80
## 4	0	8.14	0	0.538	5.950	82.0	3.9900	4	307	21.0	232.60	27.71
## 5	0	5.96	0	0.499	5.850	41.5	3.9342	5	279	19.2	396.90	8.77
## 6	25	5.13	0	0.453	5.741	66.2	7.2254	8	284	19.7	395.11	13.15
## 7	25	5.13	0	0.453	5.966	93.4	6.8185	8	284	19.7	378.08	14.44
## 8	0	4.49	0	0.449	6.630	56.1	4.4377	3	247	18.5	392.30	6.53
## 9	0	4.49	0	0.449	6.121	56.8	3.7476	3	247	18.5	395.15	8.44
## 10	0	2.89	0	0.445	6.163	69.6	3.4952	2	276	18.0	391.83	11.34
## 11	0	25.65	0	0.581	5.856	97.0	1.9444	2	188	19.1	370.31	25.41
## 12	0	25.65	0	0.581	5.613	95.6	1.7572	2	188	19.1	359.29	27.26
## 13	0	21.89	0	0.624	5.637	94.7	1.9799	4	437	21.2	396.90	18.34
## 14	0	19.58	0	0.605	6.101	93.0	2.2834	5	403	14.7	240.16	9.81
## 15	0	19.58	0	0.605	5.880	97.3	2.3887	5	403	14.7	348.13	12.03
## 16	0	10.59	1	0.489	5.960	92.1	3.8771	4	277	18.6	393.25	17.27
## 17	0	6.20	0	0.504	6.552	21.4	3.3751	8	307	17.4	380.34	3.76
## 18	0	6.20	0	0.507	8.247	70.4	3.6519	8	307	17.4	378.95	3.95
## 19	22	5.86	0	0.431	6.957	6.8	8.9067	7	330	19.1	386.09	3.53
## 20	90	2.97	0	0.400	7.088	20.8	7.3073	1	285	15.3	394.72	7.85
## 21	80	1.76	0	0.385	6.230	31.5	9.0892	1	241	18.2	341.60	12.93
## 22	33	2.18	0	0.472	6.616	58.1	3.3700	7	222	18.4	393.36	8.93
## 23	0	9.90	0	0.544	6.122	52.8	2.6403	4	304	18.4	396.90	5.98
## 24	0	7.38	0	0.493	6.415	40.1	4.7211	5	287	19.6	396.90	6.12
## 25	0	7.38	0	0.493	6.312	28.9	5.4159	5	287	19.6	396.90	6.15
## 26	0	5.19	0	0.515	5.895	59.6	5.6150	5	224	20.2	394.81	10.56
## 27	80	2.01	0	0.435	6.635	29.7	8.3440	4	280	17.0	390.94	5.99
## 28	0	18.10	0	0.718	3.561	87.9	1.6132	24	666	20.2	354.70	7.12
## 29	0	18.10	1	0.631	7.016	97.5	1.2024	24	666	20.2	392.05	2.96
## 30	0	18.10	0	0.584	6.348	86.1	2.0527	24	666	20.2	83.45	17.64
## 31	0	18.10	0	0.740	5.935	87.9	1.8206	24	666	20.2	68.95	34.02
## 32	0	18.10	0	0.740	5.627	93.9	1.8172	24	666	20.2	396.90	22.88
## 33	0	18.10	0	0.740	5.818	92.4	1.8662	24	666	20.2	391.45	22.11
## 34	0	18.10	0	0.740	6.219	100.0	2.0048	24	666	20.2	395.69	16.59
## 35	0	18.10	0	0.740	5.854	96.6	1.8956	24	666	20.2	240.52	23.79
## 36	0	18.10	0	0.713	6.525	86.5	2.4358	24	666	20.2	50.92	18.13
## 37	0	18.10	0	0.713	6.376	88.4	2.5671	24	666	20.2	391.43	14.65
## 38	0	18.10	0	0.655	6.209	65.4	2.9634	24	666	20.2	396.90	13.22
## 39	0	9.69	0	0.585	5.794	70.6	2.8927	6	391	19.2	396.90	14.10
## 40	0	11.93	0	0.573	6.976	91.0	2.1675	1	273	21.0	396.90	5.64
##	medv scored_target											
## 1	34.7	0										
## 2	18.2	1										
## 3	18.4	1										
## 4	13.2	1										
## 5	21.0	0										
## 6	18.7	0										
## 7	16.0	0										
## 8	26.6	0										
## 9	22.2	0										
## 10	21.4	0										

## 11	17.3	0
## 12	15.7	0
## 13	14.3	1
## 14	25.0	1
## 15	19.1	1
## 16	21.7	0
## 17	31.5	0
## 18	48.3	1
## 19	29.6	0
## 20	32.2	0
## 21	20.1	0
## 22	28.4	0
## 23	22.1	0
## 24	25.0	0
## 25	23.0	0
## 26	18.5	1
## 27	24.5	0
## 28	27.5	1
## 29	50.0	1
## 30	14.5	1
## 31	8.4	1
## 32	12.8	1
## 33	10.5	1
## 34	18.4	1
## 35	10.8	1
## 36	14.1	1
## 37	17.7	1
## 38	21.4	1
## 39	18.3	1
## 40	23.9	0

Appendix

For full code visit:

https://github.com/raghu74us/DATA-621/blob/master/Assignment3/621_Assignment3.Rmd